

Name : Asha Belcilda

Roll : 225229104

Lab-5 : Stemming and Lemmatization on Movie Dataset

```
In [2]: from zipfile import ZipFile
import glob
import pandas as pd
import nltk
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import linear_kernel
from nltk.corpus import stopwords
import warnings
warnings.filterwarnings('ignore')
```

EXERCISE-1

```
In [19]: file_name = "movies.zip"
with ZipFile(file_name, 'r') as zip:
    zip.printdir()
```

File Name	Modified	Size
movies/	2018-01-19 08:32:38	0
movies/12 Angry Men.txt	2018-01-17 20:40:42	1007
movies/12 Years a Slave.txt	2018-01-17 20:42:50	6451
movies/4 Months, 3 Weeks and 2 Days.txt	2018-01-17 20:37:10	1151
movies/All About Eve.txt	2018-01-17 20:33:18	1346
movies/American Graffiti.txt	2018-01-17 20:44:30	3417
movies/Boyhood.txt	2018-01-17 20:27:14	1970
movies/Casablanca.txt	2018-01-17 20:26:26	1896
movies/Citizen Kane.txt	2018-01-17 20:23:56	1483
movies/Gone with the Wind.txt	2018-01-17 20:38:10	1318
movies/Hoop Dreams.txt	2018-01-17 20:34:12	7909
movies/Manchester by the Sea.txt	2018-01-17 20:40:06	3674
movies/Moonlight.txt	2018-01-17 20:31:42	2323
movies/My Left Foot.txt	2018-01-17 20:38:50	1115
movies/Pan's Labyrinth.txt	2018-01-17 20:32:18	4431
movies/Psycho.txt	2018-01-17 20:34:46	3727
movies/Ran.txt	2018-01-17 20:43:48	2207
movies/Singin' in the Rain.txt	2018-01-17 20:29:42	782
movies/Some Like It Hot.txt	2018-01-17 20:35:40	7489
movies/The Godfather.txt	2018-01-17 20:25:32	4293
movies/Three Colors Red.txt	2018-01-17 20:28:22	2892

```
In [21]: from nltk.stem import PorterStemmer
ps = PorterStemmer()
tokenizer = nltk.tokenize.WhitespaceTokenizer()
from nltk.stem import WordNetLemmatizer
lemmatizer = WordNetLemmatizer()
from nltk.stem import LancasterStemmer
ls = LancasterStemmer()
```

```
In [24]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readlines()
        print(contents)
        print("*****")
        print(" ")
```

plantation contains the best examples both of McQueen's artistry and his occasional tendency to gild the lily. In this section, he explores, with remarkable attention to behavioral nuance, a rich array of social relationships governed by the institution of slavery—not just master and slave but mistress and slave, shopkeeper and slave, and most of all, slave and slave. The well-kept mistress of a neighboring plantation owner (Alfre Woodard) tries to convince Patsey that life as a sexual prisoner has its comparative advantages, ignoring the fact that Patsey's owner both rapes her and makes her work harder than any man in the cotton fields. Later, Patsey and Solomon become close allies after she begs him to put an end to her misery by killing her, in a scene that's as painful as it is sublimely acted. McQueen is particularly adept at exploring the intense yet guarded nature of relationships among the slaves, for whom friendship—not to speak of love—is a dangerous and necessarily transitory proposition. \n', '\n', 'But in a climactic set piece in which Solomon is forced by a gun-wielding Epps to whip Patsey nearly to death, I sometimes felt smothered by McQueen's insistence on wallowing in the extremes of human anguish. McQueen has been accused in the past of aestheticizing suffering, a critique which seems fairer to make of his two previous features (Hunger and Shame) than of this one. Hunger, in which Michael Fassbender starves himself before our eyes as the Irish hunger striker Bobbly Sands. and Sh

A. How many sentences in each file?

```
In [25]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readlines()
        for row in contents:
            sent_text = nltk.sent_tokenize(row)
            print("sentence tokenize ", len(sent_text))
```

```
sentence tokenize 0
sentence tokenize 5
sentence tokenize 0
sentence tokenize 5
sentence tokenize 0
sentence tokenize 4
sentence tokenize 4

sentence tokenize 0
sentence tokenize 5
sentence tokenize 0
sentence tokenize 2
sentence tokenize 0
sentence tokenize 5
sentence tokenize 0
sentence tokenize 3
sentence tokenize 0
sentence tokenize 1
sentence tokenize 0
sentence tokenize 3
sentence tokenize 0
```

B. How many tokens in each file?

```
In [29]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readlines()
        for row1 in contents:
            words = nltk.word_tokenize(row1)
            print("word tokenize ", len(words))
```

```
word tokenize 181
word tokenize 119
word tokenize 20
word tokenize 276
word tokenize 9
word tokenize 70
word tokenize 49
word tokenize 98
word tokenize 242
word tokenize 67
word tokenize 131
word tokenize 157
word tokenize 69
word tokenize 66
word tokenize 39
word tokenize 25
word tokenize 50
word tokenize 208
word tokenize 100
word tokenize 569
```

C. How many tokens excluding stop words in each file?

```
In [30]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readlines()
        filtered_sentence = [w for w in words if not w in stop_words]
        print("stopwords ", len(filtered_sentence))
```

```
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
stopwords 365
```

D. How many unique stems (ie., stemming) in each file? (Use PorterStemmer)

```
In [31]: def port_stemSentence(sentence):
tokenizer = nltk.tokenize.WhitespaceTokenizer()
tok = tokenizer.tokenize(sentence)
filtered_sentence = [w for w in tok if not w in stop_words]
stem_sentence = []
for word in filtered_sentence:
    stem_sentence.append(ps.stem(word))
return len(stem_sentence)
```

```
In [32]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readline()
        print("porter_stemming ")
        print(port_stemSentence(contents))
```

```
porter_stemming
96
porter_stemming
83
porter_stemming
20
porter_stemming
138
porter_stemming
63
porter_stemming
64
porter_stemming
20
porter_stemming
51
porter_stemming
131
porter_stemming
27
porter_stemming
53
porter_stemming
87
porter_stemming
35
porter_stemming
93
porter_stemming
23
porter_stemming
34
porter_stemming
52
porter_stemming
38
porter_stemming
33
porter_stemming
282
```

E. How many unique stems (ie., stemming) in each file? (Use LancasterStemmer)

```
In [33]: def lan_stemSentence(sentence):  
    tokenizer = nltk.tokenize.WhitespaceTokenizer()  
    tok = tokenizer.tokenize(sentence)  
    filtered_sentence = [w for w in tok if not w in stop_words]  
    stem_sentence = []  
    for word in filtered_sentence:  
        stem_sentence.append(ls.stem(word))  
    return len(stem_sentence)
```

```
In [34]: files = [file for file in glob.glob("movies/*")]
for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readline()
        print("lancaster_stemming ")
        print(port_stemSentence(contents))
```

```
lancaster_stemming
96
lancaster_stemming
83
lancaster_stemming
20
lancaster_stemming
138
lancaster_stemming
63
lancaster_stemming
64
lancaster_stemming
20
lancaster_stemming
51
lancaster_stemming
131
lancaster_stemming
27
lancaster_stemming
53
lancaster_stemming
87
lancaster_stemming
35
lancaster_stemming
93
lancaster_stemming
23
lancaster_stemming
34
lancaster_stemming
52
lancaster_stemming
38
lancaster_stemming
33
lancaster_stemming
282
```

F. How many unique words (ie., lemmatization) in each file? (Use WordNetLemmatizer)


```
In [54]: def lemmSentence(sentence):
tokenizer = nltk.tokenize.WhitespaceTokenizer()
tok = tokenizer.tokenize(sentence)
filtered_sentence = [w for w in tok if not w in stop_words]
lemm_sentence = []
for word in filtered_sentence:
    lemm_sentence.append(lemmatizer.lemmatize(word))
return len(lemm_sentence)
```

```
In [55]: for file in files:
    with open(file, 'r', encoding='cp1252') as f:
        contents = f.readline()
        print("lemmatization ")
        print(lemmSentence(contents))
```

```
lemmatization
96
lemmatization
83
lemmatization
20
lemmatization
138
lemmatization
63
lemmatization
64
lemmatization
20
lemmatization
51
lemmatization
131
lemmatization
27
lemmatization
53
lemmatization
87
lemmatization
35
lemmatization
93
lemmatization
23
lemmatization
34
lemmatization
52
lemmatization
38
lemmatization
33
lemmatization
282
```

EXERCISE-2

Step-1 For each movie:

Tokenize terms and build list of tokens

```
In [56]: tok = []  
for file in files:  
    with open(file, 'r', encoding='cp1252') as f:  
        contents = f.read()  
        let=tokenizer.tokenize(contents)  
        tok.append(let)  
tok
```

```
'a',  
'murder',  
'trial',  
'one',  
"man's",  
'doubts',  
'about',  
'the',  
"accused's",  
'guilt',  
'gradually',  
'overcome',  
'the',  
'rather',  
'less-than-democratic',  
'prejudices',  
'of',  
'the',  
'other',  
'eleven',
```

Find lemmatized words from the tokens

```
In [58]: tok_lem = []
for i in tok:
    for j in i:
        to_lem = lemmatizer.lemmatize(j)
        tok_lem.append(to_lem)
tok_lem

'liberalism,',
'give',
'a',
'nicely',
'underplayed',
'performance,',
'while',
'Cobb,',
'Marshall',
'and',
'Begley',
'in',

'particular',
'are',
'highly',
'effective',
'in',
'support.',
'But',
```

Step-2

Build Term-Document matrix using TfidfVectorizer

```
In [59]: for file in files:
        with open(file, 'r', encoding='cp1252') as f:
            contents = f.read()
            tok = tokenizer.tokenize(contents)
            filtered_sentence = [w for w in tok if not w in stop_words]
            tfidf = TfidfVectorizer(min_df=2, max_df=0.5, ngram_range=(1, 2))
            features = tfidf.fit_transform(filtered_sentence)
            df = pd.DataFrame(features.todense(), columns=tfidf.get_feature_names())
            print(df)
```

```
..    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...
642  ...  0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
643  ...  0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
644  ...  0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
645  ...  0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
646  ...  0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0

[647 rows x 83 columns]
      abortion  black  communist  disturbing  drama  everyone  family  gabita
\
0           0.0    0.0         0.0         0.0    1.0        0.0    0.0    0.0
1           0.0    1.0         0.0         0.0    0.0        0.0    0.0    0.0
2           0.0    0.0         0.0         0.0    0.0        0.0    0.0    0.0
3           1.0    0.0         0.0         0.0    0.0        0.0    0.0    0.0
4           0.0    0.0         1.0         0.0    0.0        0.0    0.0    0.0
..          ...    ...         ...         ...    ...        ...    ...    ...
122         0.0    0.0         0.0         0.0    0.0        0.0    0.0    0.0
123         0.0    0.0         0.0         0.0    0.0        0.0    0.0    0.0
124         0.0    0.0         0.0         0.0    0.0        0.0    0.0    0.0
```

Step-3

Take vectors of any two movies and compute cosine similarity

```
In [60]: with open(files[5], 'r', encoding='cp1252') as f:
    contents = f.read()
    tok = tokenizer.tokenize(contents)
    filtered_sentence = [w for w in tok if not w in stop_words]
    tfidf = TfidfVectorizer(min_df=2, max_df=0.5, ngram_range=(1, 2))
    movie1 = tfidf.fit_transform(filtered_sentence)
    print(movie1)
```

```
(1, 10)      1.0
(5, 2)       1.0
(12, 13)     1.0
(15, 5)      1.0
(18, 10)     1.0
(31, 20)     1.0
(35, 12)     1.0
(37, 3)      1.0
(38, 9)      1.0
(45, 10)     1.0
(46, 11)     1.0
(48, 19)     1.0
(49, 16)     1.0
(53, 8)      1.0
(54, 4)      1.0
(56, 19)     1.0
(62, 20)     1.0
(65, 12)     1.0
(69, 7)      1.0
(72, 18)     0.5773502691896258
(72, 14)     0.5773502691896258
(72, 17)     0.5773502691896258
(77, 6)      1.0
(78, 18)     0.5773502691896258
(78, 14)     0.5773502691896258
:           :
(108, 7)     1.0
(118, 5)     1.0
(121, 13)    1.0
(124, 12)    1.0
(128, 6)     1.0
(134, 10)    1.0
(138, 15)    1.0
(143, 15)    1.0
(148, 7)     1.0
(152, 1)     1.0
(154, 1)     1.0
(156, 1)     1.0
(165, 9)     1.0
(166, 0)     1.0
(172, 4)     1.0
(173, 2)     1.0
(174, 8)     1.0
(177, 10)    1.0
(179, 3)     1.0
(180, 0)     1.0
(188, 20)    1.0
(193, 7)     1.0
```

(194, 11)	1.0
(196, 12)	1.0
(203, 10)	1.0

```
In [61]: with open(files[10], 'r', encoding='cp1252') as f:
          contents = f.read()
          tok = tokenizer.tokenize(contents)
          filtered_sentence = [w for w in tok if not w in stop_words]
          tfidf = TfidfVectorizer(min_df=2, max_df=0.5, ngram_range=(1, 2))
          movie2 = tfidf.fit_transform(filtered_sentence)
          print(movie2)
```

```
(0, 15)      1.0
(1, 27)      1.0
(2, 34)      1.0
(3, 6)       1.0
(4, 8)       1.0
(7, 26)      1.0
(11, 22)     1.0
(13, 19)     1.0
(15, 20)     1.0
(17, 0)      1.0
(29, 11)     1.0
(34, 16)     1.0
(46, 35)     1.0
(52, 43)     1.0
(53, 20)     1.0
(62, 11)     1.0
(66, 20)     1.0
(67, 10)     1.0
(71, 14)     1.0
(73, 2)      1.0
(74, 18)     1.0
(77, 37)     1.0
(78, 12)     1.0
(81, 39)     1.0
(82, 20)     1.0
:           :
(323, 34)    1.0
(324, 25)    1.0
(331, 42)    1.0
(332, 19)    1.0
(333, 40)    1.0
(336, 23)    1.0
(337, 29)    1.0
(342, 31)    1.0
(343, 33)    1.0
(345, 38)    1.0
(353, 3)     1.0
(354, 11)    1.0
(356, 24)    1.0
(359, 28)    1.0
(361, 27)    1.0
(362, 34)    1.0
(366, 43)    1.0
(369, 22)    1.0
(371, 30)    1.0
(373, 41)    1.0
(379, 4)     1.0
(381, 36)    1.0
```

```
(383, 7)      1.0  
(384, 39)     1.0  
(385, 4)      1.0
```

```
In [62]: doc1 = movie1[0:10]  
doc2 = movie1[:]  
score = linear_kernel(doc1,doc2)  
print(score)
```

```
[[0. 0. 0. ... 0. 0. 0.]  
 [0. 1. 0. ... 0. 0. 0.]  
 [0. 0. 0. ... 0. 0. 0.]  
 ...  
 [0. 0. 0. ... 0. 0. 0.]  
 [0. 0. 0. ... 0. 0. 0.]  
 [0. 0. 0. ... 0. 0. 0.]]
```

```
In [ ]:
```